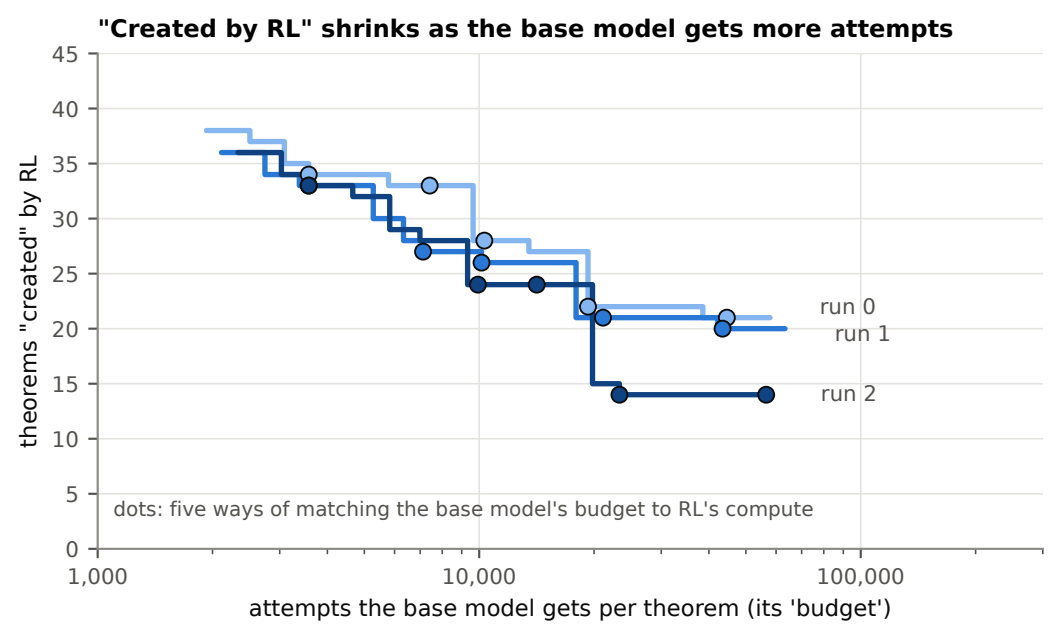
"Created by RL" shrinks as the base model gets more attempts
theorems "created" by RL
45
40
35
30
25
20
15
10
5
0
1,000
10,000
100,000
attempts the base model gets per theorem (its 'budget')
run 0
run 1
run 2
dots: five ways of matching the base model's budget to RL's compute